


EE655

ASSIGNMENT 1

Q1) Given the four key points of the mouth shown below, design at least three features that can help us monitor whether a person is smiling. **[2 Marks]**



```
import numpy as np

# WIDTH TO HEIGHT RATIO
def width_to_height_ratio(left_corner, right_corner, top_lip, bottom_lip):
    width = np.linalg.norm(np.array(left_corner) - np.array(right_corner))
    height = np.linalg.norm(np.array(top_lip) - np.array(bottom_lip))
    return width / max(height, 1) # Avoid division by zero

# LIP DISTANCE
def lip_distance(top_lip, bottom_lip):
    return np.linalg.norm(np.array(top_lip) - np.array(bottom_lip))

# CORNER DISTANCE RATIO
def corner_distance_ratio(left_corner, right_corner):
    width = np.linalg.norm(np.array(left_corner) - np.array(right_corner))
    height = abs(left_corner[1] - right_corner[1])
    return height / max(width, 1)
```

Q2) Implement a modified LeNet architecture from scratch and train it on the MNIST dataset. Your LeNet architecture must incorporate the following changes: **[3 marks]**

- Include a softmax layer at the end.
- Use $x * \text{sigmoid}(x)$ as the activation function.
- Replace average pooling with max pooling.
- Use only 3×3 filters in convolutional layers.

```
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms

# Data transformations
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

# Load MNIST dataset
train_dataset = torchvision.datasets.MNIST(root='./data', train=True, download=True, transform=transform)
train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=64, shuffle=True)

test_dataset = torchvision.datasets.MNIST(root='./data', train=False, download=True, transform=transform)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=64, shuffle=False)

# Swish-like activation function
def sigmoid_activation(x):
    return x * torch.sigmoid(x)

# Modified LeNet Model
class ModifiedLeNet(nn.Module):
    def __init__(self):
        super(ModifiedLeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=3)

        self.conv3 = nn.Conv2d(16, 120, kernel_size=3)
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(120 * 4 * 4, 84)
        self.fc2 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(sigmoid_activation(self.conv1(x)))
        x = self.pool(sigmoid_activation(self.conv2(x)))
        x = sigmoid_activation(self.conv3(x))
        x = torch.flatten(x, start_dim=1)
        x = sigmoid_activation(self.fc1(x))
        x = self.fc2(x) # No Softmax (CrossEntropyLoss applies it internally)
        return x

# Device configuration
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = ModifiedLeNet().to(device)
```

```

# Loss function & Optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Training Loop
epochs = 10
for epoch in range(1, epochs + 1):
    model.train()
    running_loss = 0.0

    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    avg_loss = running_loss / len(train_loader)
    print(f"Epoch {epoch}/{epochs}, Loss: {avg_loss:.4f}")

# Evaluation
model.eval()
correct = 0
total = 0

with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        predicted = outputs.argmax(dim=1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total
print(f'Accuracy: {accuracy:.2f}%')

```

OUTPUT:

```

Epoch 1/10, Loss: 0.1944
Epoch 2/10, Loss: 0.0502
Epoch 3/10, Loss: 0.0359
Epoch 4/10, Loss: 0.0271
Epoch 5/10, Loss: 0.0214
Epoch 6/10, Loss: 0.0167
Epoch 7/10, Loss: 0.0141
Epoch 8/10, Loss: 0.0129
Epoch 9/10, Loss: 0.0099
Epoch 10/10, Loss: 0.0079
Accuracy: 99.24%

```

Q3) Implement a modified Histogram of Oriented Gradients (HoG) feature extraction algorithm from scratch, using the "[Robert cross edge detector](#)" for computing derivatives. Extract features from images in the [Cat and Dog dataset](#) and train a Random Forest classifier. **[2 marks]**

```
import cv2
import os
import numpy as np
import matplotlib.pyplot as plt
from skimage.feature import hog
from sklearn.metrics import classification_report, confusion_matrix
from tqdm import tqdm

CATEGORIES = ["cats", "dogs"]

def load_data(folder, img_size):
    data, labels = [], []
    for category in CATEGORIES:
        path = os.path.join(folder, category)
        label = CATEGORIES.index(category)
        for img_name in tqdm(os.listdir(path), desc=f"Loading {category}"):
            img = cv2.imread(os.path.join(path, img_name))
            if img is not None:
                img = cv2.resize(img, (img_size, img_size))
                img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
                data.append(img.astype(np.float32) / 255.0)
                labels.append(label)
    return np.array(data).reshape(-1, img_size, img_size, 1), np.array(labels)

train_path = r"C:\Users\Ayush\Desktop\training_set"
test_path = r"C:\Users\Ayush\Desktop\test_set"
X_train, y_train = load_data(train_path, 128)
X_test, y_test = load_data(test_path, 128)
```

```
from scipy.ndimage import convolve

roberts_x = np.array([[1, 0], [0, -1]], dtype=np.float32)
roberts_y = np.array([[0, 1], [-1, 0]], dtype=np.float32)

def roberts_gradient(image):
    Gx, Gy = convolve(image, roberts_x), convolve(image, roberts_y)
    mag = np.sqrt(Gx**2 + Gy**2)
    orient = np.arctan2(Gy, Gx) * (180 / np.pi)
    orient[orient < 0] += 180
    return mag, orient

def batch_gradients(images):
    mags, oriens = [], []
    for img in images:
        mag, orient = roberts_gradient(img.squeeze())
        mags.append(mag)
        oriens.append(orient)
    return np.array(mags), np.array(oriens)

gradient_magnitudes, gradient_orientations = batch_gradients(X_train)
```

```

from RCedge import roberts_gradient

def hog_features(image, cell_size=8, block_size=2, bins=9):
    h, w = image.shape
    mag, orient = roberts_gradient(image)

    cells_x, cells_y = w // cell_size, h // cell_size
    hog_desc = np.zeros((cells_y, cells_x, bins))

    bin_width = 180 / bins
    for i in range(cells_y):
        for j in range(cells_x):
            cell_mag = mag[i * cell_size:(i + 1) * cell_size, j * cell_size:(j + 1) * cell_size]
            cell_orient = orient[i * cell_size:(i + 1) * cell_size, j * cell_size:(j + 1) * cell_size]
            hist = np.zeros(bins)
            for x in range(cell_size):
                for y in range(cell_size):
                    hist[int(cell_orient[x, y] // bin_width) % bins] += cell_mag[x, y]
            hog_desc[i, j, :] = hist

    blocks_x, blocks_y = cells_x - block_size + 1, cells_y - block_size + 1
    norm_hog = np.zeros((blocks_y, blocks_x, bins * (block_size ** 2)))

    for i in range(blocks_y):
        for j in range(blocks_x):
            block = hog_desc[i:i + block_size, j:j + block_size, :].flatten()
            norm_hog[i, j, :] = block / (np.linalg.norm(block) + 1e-6)

    return norm_hog.flatten()

```

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

X_train_hog = np.array([hog_features(img.squeeze()) for img in X_train])
X_test_hog = np.array([hog_features(img.squeeze()) for img in X_test])

rf_clf = RandomForestClassifier(n_estimators=100, random_state=42)
rf_clf.fit(X_train_hog, y_train)

y_pred_rf = rf_clf.predict(X_test_hog)

print("\n", classification_report(y_test, y_pred_rf))
print("\n", confusion_matrix(y_test, y_pred_rf))

```

OUTPUT:


	precision	recall	f1-score	support
0	0.68	0.76	0.72	1011
1	0.73	0.63	0.68	1012
accuracy			0.70	2023
macro avg	0.70	0.70	0.70	2023
weighted avg	0.70	0.70	0.70	2023

```

[[772 239]
 [370 642]]

```

Q4) Develop an algorithm from scratch to programmatically count the number of objects present in the binary image shown below. **[3 marks]**



```
import cv2
import numpy as np

image_path = r"C:\Users\Ayush\OneDrive\Desktop\Objects.png"
image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)

if image is None:
    raise FileNotFoundError(f"Error: Could not load image at {image_path}")

def count_objects(image):
    height, width = image.shape
    visited = np.zeros((height, width), dtype=bool)
    directions = [(-1, 0), (1, 0), (0, -1), (0, 1), (-1, -1), (-1, 1), (1, -1), (1, 1)]

    def dfs(x, y):
        stack = [(x, y)]
        while stack:
            cx, cy = stack.pop()
            if 0 <= cx < height and 0 <= cy < width and image[cx, cy] == 255 and not visited[cx, cy]:
                visited[cx, cy] = True
                stack.extend((cx + dx, cy + dy) for dx, dy in directions)

    _, binary_image = cv2.threshold(image, 128, 255, cv2.THRESH_BINARY)
    object_count = sum(dfs(i, j) or 1 for i in range(height) for j in range(width) if binary_image[i, j] == 255 and not visited[i, j])

    return object_count

num_objects = count_objects(image)
print(f"Number of objects detected: {num_objects}")
```

OUTPUT:

```
Number of objects detected: 9
```