

HOMEWORK 1 EE655 CV/DL

Ayush Kiran Badgujar
230259

January 25, 2025

Aim

The objective of this assignment is to apply the Histogram of Oriented Gradients (HOG) feature extraction algorithm to the MNIST dataset of handwritten digits and develop an image classification model. The goal is to experiment with various machine learning algorithms and achieve the best possible accuracy by tuning hyperparameters and optimizing the HOG algorithm parameters.

Theory

Histogram of Oriented Gradients (HOG)

The HOG feature extraction technique is used to detect object features in images by calculating the gradient orientation and magnitude. HOG divides an image into small spatial regions (cells), computes a histogram of gradient directions for each cell, and normalizes these histograms over overlapping regions called blocks. The HOG descriptor provides a robust representation of local patterns.

MNIST Dataset

The MNIST dataset is a well-known collection of grayscale images of handwritten digits (0-9). The training set consists of 60,000 images, and the testing set contains 10,000 images. Each image is 28x28 pixels in size.

Methodology

1. Dataset Preparation:

The dataset was downloaded from the MNIST-JPG repository. The images were organized into folders labeled by the corresponding digits (0-9).

2. **Preprocessing:**

The images were resized to 28x28 pixels (if needed). Grayscale images were used directly for HOG feature extraction.

3. **HOG Feature Extraction:**

- Cell size: 7×7
- Block size: 14×14
- Block stride: 7×7
- Number of orientation bins: 9

4. **Model Training and Evaluation:**

The following machine learning models were trained:

- Logistic Regression
- Support Vector Machine (SVM)
- Random Forest Classifier
- Naive Bayes
- k-Nearest Neighbors (k-NN)
- Neural Network (MLP)

The performance of each model was measured using accuracy and classification reports.

Code Implementation

The following Python code was used to perform the task:

```
1 import os
2 import cv2
3 import numpy as np
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.svm import SVC
6 from sklearn.ensemble import RandomForestClassifier
7 from sklearn.naive_bayes import GaussianNB
8 from sklearn.neighbors import KNeighborsClassifier
9 from sklearn.neural_network import MLPClassifier
10 from sklearn.metrics import accuracy_score, classification_report
11
12 # HOG feature extraction function
13 def extract_hog_features(images):
14     hog = cv2.HOGDescriptor(_winSize=(28, 28),
15                             _blockSize=(14, 14),
16                             _blockStride=(7, 7),
17                             _cellSize=(7, 7),
18                             _nbins=9)
19
20     features = []
21     for img in images:
22         img_resized = cv2.resize(img, (28, 28)) # Ensure the image is 28
23         hog_features = hog.compute(img_resized)
24         features.append(hog_features.flatten())
25     return np.array(features)
26
27 def load_images_from_folders(base_path):
28     images = []
29     labels = []
30     for label in os.listdir(base_path):
31         label_path = os.path.join(base_path, label)
32         if os.path.isdir(label_path):
33             for img_file in os.listdir(label_path):
34                 img_path = os.path.join(label_path, img_file)
35                 img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
36                 if img is not None:
37                     images.append(img)
38                     labels.append(int(label))
39     return np.array(images), np.array(labels)
40
41 def main():
42     train_path = "D:\\MNIST Dataset JPG format\\MNIST - JPG - training" #
43     # Replace with the path to your training folders
44     test_path = "D:\\MNIST Dataset JPG format\\MNIST - JPG - testing" #
45     # Replace with the path to your testing folders
46
47     print("Loading training dataset...")
48     train_images, train_labels = load_images_from_folders(train_path)
49
50     print("Loading testing dataset...")
```

```

48     test_images, test_labels = load_images_from_folders(test_path)
49
50     print("Extracting HOG features...")
51     X_train = extract_hog_features(train_images)
52     X_test = extract_hog_features(test_images)
53
54     y_train = train_labels
55     y_test = test_labels
56
57     accuracies = {}
58
59     # Logistic Regression
60     print("Training Logistic Regression model...")
61     lr_model = LogisticRegression(max_iter=1000)
62     lr_model.fit(X_train, y_train)
63     lr_predictions = lr_model.predict(X_test)
64     accuracies['Logistic Regression'] = accuracy_score(y_test,
65                                                         lr_predictions)
66
67     # Support Vector Machine
68     print("Training Support Vector Machine model...")
69     svm_model = SVC(kernel='linear', C=1.0)
70     svm_model.fit(X_train, y_train)
71     svm_predictions = svm_model.predict(X_test)
72     accuracies['SVM'] = accuracy_score(y_test, svm_predictions)
73
74     # Random Forest
75     print("Training Random Forest model...")
76     rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
77     rf_model.fit(X_train, y_train)
78     rf_predictions = rf_model.predict(X_test)
79     accuracies['Random Forest'] = accuracy_score(y_test, rf_predictions)
80
81     # Naive Bayes
82     print("Training Naive Bayes model...")
83     nb_model = GaussianNB()
84     nb_model.fit(X_train, y_train)
85     nb_predictions = nb_model.predict(X_test)
86     accuracies['Naive Bayes'] = accuracy_score(y_test, nb_predictions)
87
88     # k-Nearest Neighbors
89     print("Training k-Nearest Neighbors model...")
90     knn_model = KNeighborsClassifier(n_neighbors=3)
91     knn_model.fit(X_train, y_train)
92     knn_predictions = knn_model.predict(X_test)
93     accuracies['k-NN'] = accuracy_score(y_test, knn_predictions)
94
95     # Neural Network
96     print("Training Neural Network model...")
97     nn_model = MLPClassifier(hidden_layer_sizes=(100,), max_iter=300,
98                               random_state=42)
99     nn_model.fit(X_train, y_train)
100    nn_predictions = nn_model.predict(X_test)
101    accuracies['Neural Network'] = accuracy_score(y_test, nn_predictions)

```

```

100
101 # Find the best model
102 best_model_name = max(accuracies, key=accuracies.get)
103 print(f"\nBest Model: {best_model_name} with accuracy {accuracies[
    best_model_name]}")
104
105 # Generate the classification report for the best model
106 if best_model_name == "Logistic Regression":
107     best_predictions = lr_predictions
108 elif best_model_name == "SVM":
109     best_predictions = svm_predictions
110 elif best_model_name == "Random Forest":
111     best_predictions = rf_predictions
112 elif best_model_name == "Naive Bayes":
113     best_predictions = nb_predictions
114 elif best_model_name == "k-NN":
115     best_predictions = knn_predictions
116 elif best_model_name == "Neural Network":
117     best_predictions = nn_predictions
118
119 print(f"\nClassification Report for {best_model_name}:")
120 print(classification_report(y_test, best_predictions))
121
122 if __name__ == "__main__":
123     main()

```

Listing 1: Python Code for MNIST Classification with HOG Features

Results and Output

The following table summarizes the model performance:

Table 1: Model Performance Summary

Model	Description	Accuracy (%)
Logistic Regression	Linear model for classification	95.72
Support Vector Machine	Linear SVM kernel	99.00
Random Forest	Ensemble of decision trees	97.85
Naive Bayes	Probabilistic model	83.49
k-Nearest Neighbors (kNN)	Distance-based classification	95.31
Neural Network (MLP)	1-hidden layer MLP model	97.12

Classification Report for Best Model (SVM)

Table 2: Classification Report for Support Vector Machine (SVM)

Class	Precision	Recall	F1-Score	Support
0	0.99	1.00	0.99	980
1	0.99	1.00	0.99	1135
2	0.99	0.99	0.99	1032
3	0.99	0.99	0.99	1010
4	0.99	0.99	0.99	982
5	0.99	0.99	0.99	892
6	0.99	0.99	0.99	958
7	0.99	0.98	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.98	0.98	1009
Accuracy			0.99	10000
Macro avg	0.99	0.99	0.99	10000
Weighted avg	0.99	0.99	0.99	10000

Effect of HOG Parameters on Results

The following observations were made by varying HOG parameters:

- Cell Size: Decreasing the cell size (e.g., 4×4) increased the feature dimensions, leading to a higher computational cost without significant accuracy improvement. Larger cell sizes (e.g., 8×8) resulted in loss of fine details and reduced accuracy.
- Block Size: Smaller block sizes (e.g., 10×10) improved accuracy slightly by better capturing local patterns but increased computational complexity.
- Block Stride: Larger strides (e.g., 10×10) reduced feature overlap and decreased accuracy, while smaller strides improved overlap and boosted performance.
- Number of Orientation Bins: Increasing bins (e.g., 12 or 18) captured finer details, but bins larger than 12 showed diminishing returns.

Conclusion

The Support Vector Machine (SVM) demonstrated the best performance in classifying hand-written digits from the MNIST dataset, achieving an accuracy of 99%. Experimentation with HOG parameters revealed a trade-off between computational efficiency and accuracy. Smaller cell sizes and finer orientation bins improve accuracy but at the cost of higher computational complexity.